

Demo & Intro

0:00

hey what's up guys I have something that will absolutely blow your mind this is

0:05

the computer playing Super Mario Bros and guess what the computer has no knowledge about the internal State or

0:11

the mechanics of the video game rather it's simply looking at the frames and then deciding what buttons to press

0:18

exactly how humans do it impressive right well if you want to learn how to

0:24

implement this stay tuned because that's exactly what I'll be sharing with you guys today and I cannot wait

0:30

this was done using a subclass of machine learning algorithms called reinforcement learning reinforcement

0:36

learning is the coolest thing ever in fact it's been used to train drones and

0:41

robots find more efficient ways for matrix multiplication and even make the training process for other machine

0:48

learning algorithms faster oh and not to mention it was also used to train

0:53

everyone's favorite chatbot chat GPT reinforcement learning is pretty different from its cousin's supervised

0:59

and unsupervised learning in supervised learning the model is given a data set

1:04

where each point has a correct answer associated with it whether it's categorical or numerical the model's job

1:11

is then to predict these values accurately an unsupervised learning the data points

1:17

don't have a correct answer but rather the models tasked with extracting General Trends or insights reinforcement

1:24

learning is its own Beast the objective of reinforcement learning is for the agent to maximize its rewards inside of

1:31

the environment we give the agent a positive reward whenever it exhibits behaviors we want

1:36

to persist and negative ones for behaviors that we don't want it's kind of like teaching a dog new tricks

1:42

another thing to note is that in supervised and in unsupervised learning you must have the data pre-collected

1:49

before you feed it to the model whereas in reinforcement learning we'll see today our agent go out into the world

1:55

collect its own experiences and then learn from it it's a completely different Paradigm when it comes to

2:01

training I wanted to share an implementation of the ddqn algorithm or the double deep Q Network I'm not going

2:09

to dive into the nitty-gritty of reinforcement learning theory I'd rather just share a more intuitive

2:14

understanding for those that have a base level understanding of machine learning but haven't really been exposed to

2:19

reinforcement learning before however for those that are interested I'm going to leave some of my favorite

2:25

reinforcement learning resources down in the description below along with the paper that originally pioneered the

2:31

double deep Q Network algorithm I first implemented this algorithm when I was a sophomore in college and I still

2:37

remember the excitement that I felt when I first watched the model start to train learn and eventually beat the level

2:44

I was absolutely awestruck I mean I understood all of the math and each line of code that I wrote to implement the

2:50

algorithm however there's just something that's still magical about reinforcement learning that gets me excited every

2:56

single time and my goal today is to be able to package that excitement and share it with you guys through this

Key Reinforcement Learning Vocabulary

3:02

video let's start off with the basics in order to understand the fundamentals of reinforcement learning we need to Define

3:08

these following terms let's understand each of them in the context of Super Mario Bros first is the agent the agent

3:17

today will be our neural network that's controlling Mario it's the one responsible for making decisions and

3:23

taking the optimal action the environment is what our agent will be interacting with it's going to be level

3:30

one one of Super Mario Bros on the NES a state is a snapshot of the environment

3:36

at any given time it's essentially what our agent sees today we'll be using four

3:41

consecutive frames from the game the reason we're using consecutive frames is so our agent can see motion for example

3:49

is the Goomba coming towards us or away from us well for us to know we need to

3:54

know the previous frames oh well hopefully our agent knows to jump actions are the inputs that the

4:01

agent has available to it to give to the environment today the agent will have access to these five button combinations

4:09

to press first one is with no buttons pressed and then the agent is able to move right a is used to jump and the

4:16

longer it's held the higher Mario will jump and then B is used to move faster

4:22

once the environment gives the agent a state and the agent responds with an

4:27

action the environment also gives the agent a reward a reward essentially

4:33

helps the agent determine whether or not the action it took inside of that state was good or bad today we'll be giving

4:40

the agent a plus one reward for each unit to the right that it makes in the video game it'll receive a negative one

4:46

reward for each second that passes by in the games clock this is to disincentivize the agent from standing

4:52

still and then the agent will get a massive negative 15 points if it dies

4:59

now the concept of a negative reward may seem foreign to some however remember the goal of the rewards is to guide the

5:06

agent to Optimal actions and negative rewards still help the agent do just

5:11

that everything that we've talked about so far is summarized in this diagram it

5:17

essentially outlines how the agent and the environment interact with one another the agent gives the environment

5:23

actions and the environment returns States and rewards an episode for us

5:29

today will just be one attempt at the level an episode will end whether Mario dies reaches the flag or Runs Out of

5:37

Time a policy is a function that takes in a state and returns an action it's

5:43

what our agent will use to make its decisions some reinforcement learning algorithms will have the policy be a

5:50

probability distribution however today in our ddqn algorithm implementation

5:56

we'll be using the Epsilon greedy approach more on that later a value function takes in a state and returns

6:03

how valuable that state is in reinforcement learning it's natural to

6:08

view some states as more valuable than others for example the state where you're right next to the flag is

6:15

probably a little bit more valuable than the one where you're falling down a hole we're not going to be working directly

6:20

with the value function today however I thought I'd mention it because a lot of other reinforcement learning literature

6:26

brings it up an action value function takes in a state and an action and then

6:33

returns how valuable this pair is as you can see on the screen this state action

6:39

pair is probably more valuable than this state action pair I'm sure you can see why approximating this function

6:44

accurately today is our goal because it's going to let our agent take the

6:49

action inside of in state that has the highest value alright quick review we're

6:55

using reinforcement learning to train the computer to play Super Mario Bros our agent or the neural network that's

7:01

controlling Mario will receive a state from the environment remember a state is

7:06

four consecutive frames from the video game the agent will then take an action

7:11

if the action was good we'll give it a positive reward if the action was bad it will give it a negative reward and then

7:18

the agent will use these rewards to train itself over time time to take better and better actions eventually

7:25

we're hoping that the agent will learn how to beat the level and reach the flag that's waiting for it at the end that

7:32

was a lot to take in thank you guys for bearing with me we have three more Concepts to cover first is the Epsilon

7:37

greedy approach second is the replay buffer and third are some more details about the action value function after

7:44

that we'll assemble the algorithm and then implement it in code the Epsilon greedy approach is the strategy our

Epsilon-Greedy Approach

7:50

agent will be using today to choose its actions however to understand Epsilon greedy we first must take a step back

7:57

and understand the explore exploit dilemma a problem that's at the heart of reinforcement learning let's say our

8:04

agent has learned that jumping over Koopas is the most effective way to not die good I mean I guess it gets the job

8:11

done right however as I'm sure my fellow Gamers know that jumping on a Koopa

8:17

actually turns its shell into a weapon that can be used to eliminate other enemies if our agent constantly exploits

8:24

or takes the action that it deems to be the best in the current moment then it will never get a chance to explore other

8:30

potentially Superior strategies so how do we navigate this explore exploit dilemma enter the world of the Epsilon

8:36

greedy approach a clever solution to the explore exploit dilemma with probability Epsilon our agent Ventures into the

8:43

unknown taking random actions and with the remaining probability $1 - \epsilon$ it'll take the action that it's

8:49

confident is the best one it's essential to note that this best action of the agent is dependent only on what it knows

8:55

in the current moment it's not necessarily the best action overall initially when our agent knows nothing

9:01

about the environment we'll set ϵ to 1 ensuring maximum exploration and

9:06

then over time as our agent continues to learn about the environment and its Dynamics will start to taper off ϵ

9:13

from one to a very small non-zero number we'll always keep ϵ greater than zero just so our agent is always on the

9:20

lookout for new strategies this means that if our agent did learn to jump over Koopas we're hoping that with

9:26

the ϵ -greedy approach it'll eventually randomly jump on one and then realize the benefits of doing so now

Replay Buffer

9:33

let's cover the replay buffer it's essentially a storage for our past experiences and rewards it'll be used to

9:39

train the neural network there are a couple of reasons why we use a replay buffer rather than training on the

9:44

experiences as we collect them the first is because sequential experiences have a lot of correlation which can lead to a

9:52

lot of instability when training the model to eliminate this instability we randomly sample from our replay buffer

9:58

and the second reason is that we're now able to reuse data which actually improves data efficiency in the replay

10:05

buffer we'll be storing tuples of the state action taken reward received next

10:11

state and a Boolean flag indicating if the episode is done for any reason we'll see how each of these come into play

10:17

when we sample the replay buffer to train our Network now let's pull up the action value function that we mentioned

Action-Value Function Intuition

10:23

earlier this is often referred to as the Bellman equation remember our primary objective today is to approximate this

10:30

function enabling our agent to make optimal decisions in any given State the goal isn't just about seeking immediate

10:36

gratification instead the agent aims to maximize its rewards over the long run humans intuitively do this as well for

10:44

example we may take a slightly longer route with less coins if it leads us to an item box at the end that contains the

10:50

invincibility star math notation is scary so let me translate this equation to English the value of taking action a

10:58

in state s is equal to the reward you get in the next time step plus the value

11:03

of taking the best action in the next state S' multiplied by a discount

11:09

Factor the discount Factor makes rewards from future States less valuable in

11:14

reinforcement learning this is important because future rewards are less predictable than the current one due to

11:20

the environment's stochastic in the Stanford marshmallow experiment children were given a marshmallow and

11:26

were told that they could either eat it right now or wait 15 minutes for a second marshmallow it was to measure

11:33

their self-control when it came to delayed gratification essentially the researchers were measuring the

11:39

children's γ or their discount Factor one marshmallow now or two in the

11:44

future our agent is kind of in a similar situation except for the children their

11:50

second marshmallow was guaranteed unfortunately our agent does not have that luxury as such it's ideal for it to

11:58

value a guaranteed marshmallow in the present more than a potential marshmallow in the future a γ of one

12:05

means that there is no discounting while a γ of zero means the agent is completely myopic we want neither but

12:13

rather to strike a balance between the two this action value function is recursive in nature this means that the

12:20

value of a specific State action pair is not just based on the immediate reward rather it also depends on the value of

12:27

the subsequent State action pair dive deeper and you'll see that this subsequent value in turn depends on its

12:33

immediate reward and the value of the next state action pair and the chain continues we assume that we take the

12:39

optimal action the action with the highest Q value at each recursive level thus the max function when you unravel

12:46

this recursive structure you'll find that the value of any state action pair is essentially the sum of its immediate

12:52

reward and the accumulated discounted feature rewards until the end of the episode this first equation once you

12:59
assume you're taking the best action at each time step condenses down to this after you distribute the Gammas

13:05
appropriately the further out a reward is the more times it's multiplied by gamma still a little bit confused don't

13:12
worry I was too when I first learned this let's try to tackle this from my more visual approach take each ribbon on

13:18
the screen to be a reward the agent will get until it reaches the flag remember that these rewards are granted after

13:24
each step the agent takes in the environment and the number on the screen is arbitrary the ribbons are getting

13:30
progressively smaller to signify the discounting assuming the agent is currently in state s and takes action a

13:36
the value of this state action pair is the sum of this entire sequence of rewards discounted appropriately if S

13:44
Prime is the next state after taking action a in state s and a' is the

13:49
best possible action in S' then the value of S' a' will be this

13:54
sum again discounted appropriately we know the reward R we got from going from

13:59
 s to S' so we can use that to help tune the estimate of the current state action pair's value to train our agent

14:07
We compare its predicted value of a state action pair to the Target value remember the target value is the

14:13
immediate reward from a move plus the discounted value of the best possible next move essentially We

14:19
compare its predicted value of the sum of these ribbons to the predicted value of the sum of these ribbons plus the

14:26
ribbon whose value we know with certainty because the agent just received that reward while both the 14:32

prediction and the target are estimates the target is considered to be more reliable why is that the case well

14:39
because the target is based on fewer future assumptions as a component of the target is the reward directly observed

14:45
from the agent's recent action making it grounded in immediate experience rather

14:50
than speculative projection once we have our predicted value and our Target value we can use the following equation to

14:57

update our estimate for those of you familiar with gradient descent you might recognize this to be exactly that Alpha

15:04

is our learning rate or the size of the steps our Network's parameters take as they converge towards their optimal

15:09

values the part inside of the square brackets is the error or the difference between our predicted and Target values

15:15

this is exactly the derivative of the mean squared error loss function I'm sorry if all of that was overwhelming I

The DDQN Algorithm

15:21

know it was for me the very first time I learned it however the intuition of the action value function is the hardest

15:27

part I swear it gets easier from here so how do we even begin to approximate this function that has such a vast space of

15:33

potential State action pairs well we turn to our friend the neural network remember neural networks are great at

15:40

approximating functions and extremely high Dimensions the specific architecture we'll be using today is the

15:47

convolutional neural network or a CNN CNN's excel at extracting meaningful

15:52

visual features from its inputs whether its enemies holds or the flag if you're not too familiar with CNN's I highly

15:59

recommend nvidia's learning deep learning book it'll cover literally everything you need to know about CNN

16:05

the input size will be the dimensions of our state which we'll see in just a sec the number of neurons in the output

16:12

layer will be the number of actions available to us which in this case will be five the value in each neuron

16:19

represents the predicted Q value for that Associated action paired with the

16:24

input State when running this algorithm we're actually going to have two identical copies of our neural network

16:30

the first one will be the online Network this is the one that will actually be training the second one will be the

16:37

target Network the one we'll be using for the ground Truth for our predictions to give to our loss function the target

16:43

Network won't be trained however we'll be intermittently copying the weights from the online network over to the

16:49

Target Network this is to increase stability during the training process to train the online Network we sample

16:56

the replay buffer to get a random Tuple containing a state action reward next state and done flag we then pass in the

17:04

state to the online Network to get the predicted Q values for the five actions the different sized controllers

17:10
represent how the network values each action differently however we only care about the action that was taken and is

17:17

present in the sampled Tuple this is our predicted value for this state action

17:22

pair to see how accurate it is we calculate our Target value to do so we pass in the

17:29

next state into our Target Network to get the predicted Q values for the five actions again this time we take the

17:37

highest predicted Q value represented by the largest controller to compute our

17:42

Target value we take the reward from our sample and add it to the predicted Q value of the next state multiplied by

17:49

gamma our discount Factor we would then pass these two values into our loss

17:54

function and perform one learning step for our online Network this technique is known as bootstrapping because our poor

18:02

little network is essentially pulling itself up by its bootstraps it's using a crappy Target Network to train our

18:09
crappy online Network and then as the online Network improves just a little bit we'll copy over the weights from the

18:15

online Network to the Target Network and now we're using a slightly less crappy Target Network to continue training our

18:22

online Network we'll repeat keep this process tens of thousands of times and eventually we'll start to see our agent

18:28

actually start to learn I know this seems very Jank like we're using a moving estimate to train another

18:34

estimate however in reality we'll see how stable it actually is so let's put

DDQN Pseudocode

18:40

together all of these moving parts and finally assemble our ddq1 algorithm we start off with initializing our starting

18:46

variables and resetting our environment which gives us the starting State second given the state we'll choose an action

18:53

with our Epsilon greedy approach third we'll use this action to take one

18:59

step inside of the environment as mentioned previously the environment will return a new state and a reward

19:05

we'll then take the current state action reward the new state and the done flag

19:13

and store them all in our replay buffer then we'll do one learning step which involves sampling the replay

buffer and

19:20

training the online Network then we'll take our next state and assign it to our

19:25

current state we'll repeat all of this until the episode is done and finally

19:31

we'll repeat this inside the loop for however many episodes we want to train for so let's finally get to coding this

Implementation in Code

19:39

up for the NES emulator we'll be using a library called gym Super Mario Bros it

19:45

handles all of the complexities of the emulator side of things for us it gives us neatly packaged python objects for

19:51

the States along with a simple way to input actions it does the this via a

19:56

simple API that's based on openai's gym API openai's gym library is a wildly

20:03

popular collection of different reinforcement learning environments that researchers or enthusiasts like us can

20:10

use to try out different algorithms they're incredibly easy to set up and use these days the most common approach

20:17

is to use gymnasium which is a maintained Fork of openai's gym Library by an outside team I wonder what openai

20:25

is working on these days also all of the code will be linked on GitHub right down in the description and if you ever need

20:32

to reference it throughout this video or even afterwards please feel free to do so so let's start off with a simple

20:38

example let's first import the libraries we'll be limiting our agent to only press these following button

20:44

combinations which is why we're importing write only then we create our environment object

20:50

and wrap it with the joypad space wrapper ensuring that the actions in write only are the only ones the

20:57

environment can accept then we set our done flag to false and

21:02

reset our environment this while loop will run for as long as the agent is still alive

21:08

first let's have our agent only press right on the d-pad we'll then take one step in the environment with that action

21:15

and get our done flag ignoring the other return values from the step function for now and finally because on line 6 we

21:24

have our render mode set to human calling env.render will display our

21:29

environment to our screen so we can see what our agent is doing as expected our

21:35

agent is not doing so well to add some dynamism let's have it randomly choose an action from our action

space

21:42

unfortunately our agent still sucks good thing we have an entire reinforcement

21:48

learning algorithm up Our Sleeve let's implement the wrappers first rappers are essentially a way that we can modify our

21:55

environment's outputs if you want some more information about writing your own wrappers I highly recommend checking out

22:01

the gymnasium documentation which is linked below we're going to be using four different

22:06

wrappers today the wrapper we're implementing will skip frames meaning it'll just take the action inputted for

22:12

the first frame and reapply it for however many frames we want to skip we'll aggregate the rewards over those

22:19

four frames this is useful for us because consecutive frames have a lot of overlap so it's redundant to reprocess

22:26

everything the next three rappers are already implemented by gymnasium resize

22:32

observation will change the dimensions of a frame from 240 by 256 pixels to 84

22:40

by 84. this is just to reduce the computational load similarly grayscale

22:46

observation will turn the frame from having a red green and blue channel to just one reducing the amount of data our

22:53

network will have to crunch finally so our Network can see motion we're going to take four consecutive

23:00

frames and stack them on top of each other to finally create our state object that will pass into our Network each

23:07

final State object comprised of a stack of four processed frames encapsulates data from 16 original game frames due to

23:15

the combined effects of frame skipping resizing and stacking in our

23:20

pre-processing steps the code for the wrappers is pretty straightforward as we're only implementing one of them we first

23:27

override the Constructor to take the number of frames to skip then we override the step function whenever step

23:34

is called we use the for Loop to take skip number of steps aggregating the

23:39

rewards the apply wrappers function applies our wrappers one by one

23:44

encapsulating our base environment like the layers of an onion at the end it

23:50

returns our fully wrapped environment object we'll use this function in our main.pi

23:56

let's next create our neural network with pi torch as mentioned previously we'll be using a convolutional

neural
24:03
network we're going to have three convolutional layers followed by two linear layers the input shape will be 4
24:10
by 84 by 84 which is the dimensions of our state the number of actions will be five the
24:16
underscore get conv out function performs a dummy forward pass through the convolutional layers to get the
24:23
number of neurons we need in our first linear layer it's 3136 with this specific architecture but
24:31
the function allows us to dynamically calculate it where we to change anything in the future
24:37
the freeze flag allows us to prevent pie Torch from calculating the gradients which we'll need for the Target Network
24:43
remember we only use the target Network to calculate the correct values we want
24:48
our online Network to predict then finally we add our forward pass and then move the network to whichever device
24:54
we're training on whether it's the CPU or GPU let's next create the agent class
25:00
these will be our Imports along with pytorch and numpy we'll also import our
25:05
neural network that we just created for the replay buffer we'll be using pi torch's built-in tensor dict replay
25:12
buffer which will hold python dictionaries with tensors as the values
25:17
the storage mechanism it'll be using is pi torch's lazy mem map storage which
25:23
will use memory mapped files for easy access to our experiences along with
25:28
alleviating RAM usage we could have just used a python list for our experiences from which we sample
25:35
but that means sampling would be really slow and we'd be using a lot of RAM
25:40
for the Constructor we'll take in the dimensions of our state and the number of actions our agent has access to we'll
25:47
also use a learn step counter which keeps track of how many times we've trained our Network
25:53
we then set up our hyper parameters first is the learning rate or Alpha this
25:59
is the size of the steps the network will take when it's updating its weights second is gamma which is our discount
26:05
Factor remember we want the agent to Discount future Rewards
26:10
third is the starting value of Epsilon which will be set to 1 initially to encourage exploration
26:16
fourth is the Epsilon Decay factor which will multiply Epsilon with after every
26:21
time step fifth is our minimum Epsilon so we always maintain some likelihood to

26:27
explore sixth is the batch size for our training and finally is how often we'll sync the
26:35
target Network weights with the online Network after the hyper parameters comes the networks again
making sure the
26:42
target Network's parameters are frozen after adding our Optimizer and loss we create our replay buffer
with the
26:49
capacity to hold a hundred thousand experiences then let's write our choose
26:55
action function that uses the Epsilon greedy approach as we can see when the random number is less
than our Epsilon
27:02
value it'll choose a random action since Epsilon starts off at 1 the action
27:07
chosen will always be random as the value of Epsilon decays the probability
27:12
of using our online Network to choose the action with the highest Q value increases also because we
want the index
27:19
itself not the estimated value of that action we're taking the ARG Max our
27:25
Decay Epsilon function multiplies the current value of Epsilon with our Decay factor and ensures it doesn't
go below
27:32
the minimum Epsilon threshold the store in memory function takes in the tensors that we want to put into
our
27:39
replay buffer organizes them in a dictionary and adds them to the buffer
27:44
the sync networks function checks if enough learning steps have passed and if so it'll copy over the
weights of the
27:51
online Network to the Target Network now for the most important function the
27:57
learn function in the agent class we first validate that there are enough experiences in our replay buffer to
28:03
sample a batch from then we call the sync networks function we then clear our
28:08
gradients followed by sampling the replay buffer we store the results in the states actions rewards next
States
28:15
and duns variables we passed the state's tensor through the online Network to get
28:22
our predicted values we index by the actions that we actually took because we only perform back
propagation on those
28:29
values then we calculate our Target values we pass in the next States through our
28:36
Target Network and get the value that the best action yields we then multiply
28:41
it by gamma and added to the rewards we got in the current state the one minus
28:46
duns part sets all future rewards to zero if we're in a terminal state

28:52
we then calculate the loss using our predicted and Target Q values performing
28:57
back propagation to calculate the gradients and then perform a step of gradient descent with those
gradients we
29:04
also then increment our learn step counter and Decay Epsilon finally let's create our main dot Pi in this
file
29:12
we'll set up our environment as we saw previously and then create the full training loop again starting off
with
29:18
our Imports this time adding our apply wrappers function and our agent class
29:24
we then create our environment applying the joypad space wrapper plus the other four that we talked
about earlier we
29:31
also create our agent object and pass in the dimensions of our state and the number of actions from our
environment
29:37
finally the training Loop looks very similar to our pseudo code for each episode we collect experiences
and learn
29:44
from them until the episode is done first we choose an action given the starting State again this will be
29:50
completely random at first then we take a step in the environment and store the state action reward new
state and done
29:58
flag in our replay buffer we then perform one learning update then we set
30:04
our current state to our new state and repeat for training for 50 000
30:09
iterations can take a while I have a RTX 3080 and it took me about two days of
30:14
training non-stop however I'm excited to see what optimizations you guys can uncover by tooting the
hyper parameters
The AI Beats the Level!
30:21
but now the moment you've all been waiting for after we finish training
30:27
this is the final result
30:39
thank you
Conclusion
30:56
well we did it guys we finally trained the computer to play Super Mario Bros I
31:02
hope you guys feel the magic of reinforcement learning that I've always felt in fact even bigger than that I
31:08
hope you realize that machine learning is so much bigger than just predicting labels from a data set or
trying to
31:14
Cluster some points anyways if you have any questions about anything that we talked about today
31:21
please leave a comment down below and until then see you guys next time

